

# EN3160 Assignment 3 on Neural Networks

November 12, 2024

**Name:** Abithan A.

**Index No.:** 210019E

[View GitHub Repository](#)

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
```

```
[2]: # 1. Dataloading

transform = transforms.Compose(
    [ transforms.ToTensor() ,
      transforms.Normalize((0.5 , 0.5, 0.5) , (0.5 , 0.5, 0.5))])

batch_size = 50

trainset = torchvision . datasets .CIFAR10(root=r'D:\semi 5\IP&CV\Assignment_
↳3\data', train=True, download=True, transform=transform)

trainloader = torch. utils .data.DataLoader(trainset , batch_size=batch_size ,
↳shuffle=True, num_workers=2)

testset = torchvision . datasets .CIFAR10(root=r'D:\semi 5\IP&CV\Assignment_
↳3\data', train=False , download=True, transform=transform)

testloader = torch. utils .data.DataLoader( testset , batch_size=batch_size ,
↳shuffle=False , num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```

[3]: # 2. Define Network Parameters
Din = 3 * 32 * 32 # Input size (flattened CIFAR=10 image size)
K = 10 # Output size (number of classes in CIFAR=10)
std = 1e-5

# Initialize weights and biases
w = torch.randn(Din, K) * std # One layer: directly map input to output
b = torch.zeros(K)

# Hyperparameters
iterations = 20
lr = 2e-6 # Learning rate
lr_decay = 0.9 # Learning rate decay
reg = 0 # Regularization
loss_history = []

[4]: # 3. Training Loop
for t in range(iterations):
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # Get inputs and labels
        inputs, labels = data
        Ntr = inputs.shape[0] # Batch size
        x_train = inputs.view(Ntr, -1) # Flatten input to (Ntr, Din)
        y_train_onehot = nn.functional.one_hot(labels, K).float() # Convert labels to one-hot encoding

        # Forward pass
        y_pred = x_train.mm(w) + b # Output layer activation

        # Loss calculation (Mean Squared Error with regularization)
        loss = (1 / Ntr) * torch.sum((y_pred - y_train_onehot) ** 2) + reg * torch.sum(w**2)
        loss_history.append(loss.item())
        running_loss += loss.item()

        # Backpropagation
        dy_pred = (2.0 / Ntr) * (y_pred - y_train_onehot)
        dw = x_train.t().mm(dy_pred) + reg * w
        db = dy_pred.sum(dim=0)

        # Parameter update
        w -= lr * dw
        b -= lr * db

    # Print loss for every epoch
    if t % 1 == 0:

```


```
print(f"Epoch {t +1}/{ iterations}, Loss:{ running_loss/len (trainloader )} " )
```

```
# Learning rate decay  
lr *= lr_decay
```

```
Epoch 1/20, Loss:0.9769474459886551  
Epoch 2/20, Loss:0.9498618963956833  
Epoch 3/20, Loss:0.9361145228743554  
Epoch 4/20, Loss:0.9275588639378548  
Epoch 5/20, Loss:0.921618634045124  
Epoch 6/20, Loss:0.9172142955064774  
Epoch 7/20, Loss:0.9138009185791016  
Epoch 8/20, Loss:0.9110727261900902  
Epoch 9/20, Loss:0.9088412716984748  
Epoch 10/20, Loss:0.9069875936508178  
Epoch 11/20, Loss:0.9054266312718391  
Epoch 12/20, Loss:0.9040986114740371  
Epoch 13/20, Loss:0.9029594156742096  
Epoch 14/20, Loss:0.9019756987094879  
Epoch 15/20, Loss:0.9011212394833564  
Epoch 16/20, Loss:0.9003760864138604  
Epoch 17/20, Loss:0.8997233058214188  
Epoch 18/20, Loss:0.8991496341824532  
Epoch 19/20, Loss:0.8986439597010613  
Epoch 20/20, Loss:0.8981972292661667
```

[5]: *# 4. Plotting the Loss History*

```
plt.plot(loss_history)  
plt.title("Loss_History")  
plt.xlabel("Iteration")  
plt.ylabel("Loss")  
plt.show()
```



```
[6]: # 5. Calculate Accuracy on Training Set
correct_train = 0
total_train = 0
with torch.no_grad():
    for data in trainloader:
        inputs, labels = data
        Ntr = inputs.shape[0]
        x_train = inputs.view(Ntr, -1)
        y_train_onehot = nn.functional.one_hot(labels, K).float()

        # Forward pass
        y_train_pred = x_train.mm(w) + b
        predicted_train = torch.argmax(y_train_pred, dim=1)
        total_train += labels.size(0)
        correct_train += (predicted_train == labels).sum().item()

train_acc = 100 * correct_train / total_train
print(f"Training accuracy :{ train_acc :.2f}%")
```

Training accuracy :32.21%

```
[7]: # 6. Calculate Accuracy on Test Set
correct_test = 0
total_test = 0

with torch.no_grad():
    for data in testloader:
        inputs, labels = data
        Nte = inputs.shape[0]
        x_test = inputs.view(Nte, -1)
        y_test_onehot = nn.functional.one_hot(labels, K).float()

        # Forward pass
        y_test_pred = x_test.mm(w) + b
        predicted_test = torch.argmax(y_test_pred, dim=1)
        total_test += labels.size(0)
        correct_test += (predicted_test == labels).sum().item()

test_acc = 100 * correct_test / total_test
print(f"Test accuracy : {test_acc : .2f}%")
```

Test accuracy : 32.41%

1. Listing 1 shows the code for a single dense layer network with manually computed forward path and backpropagations. Do the following changes

- Add a middle layer with 100 nodes and a sigmoid activation.
- Use cross-entropy loss
- Run the network for 10 epochs and report the training and test accuracies.

```
[ ]: # Model definition
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.fc1 = nn.Linear(3 * 32 * 32, 100) # Input to hidden layer with
        ↪ 100 nodes
        self.sigmoid = nn.Sigmoid() # Sigmoid activation
        self.fc2 = nn.Linear(100, 10) # 100 Hidden layer nodes to
        ↪ output layer

    def forward(self, x):
        x = self.fc1(x)
        x = self.sigmoid(x)
        x = self.fc2(x)
        return x
```

```
[ ]: model = NeuralNetwork()
criterion = nn.CrossEntropyLoss()
```

```

optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9) # learning
↳rate = 0.01, momentum = 0.9

# Dataloading and transformations
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.
↳5, 0.5, 0.5), (0.5, 0.5, 0.5))])
batch_size = 50

trainset = torchvision.datasets.CIFAR10(root=r'D:\semi 5\IP&CV\Assignment
↳3\data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
↳shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root=r'D:\semi 5\IP&CV\Assignment
↳3\data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
↳shuffle=False, num_workers=2)

```

Files already downloaded and verified

Files already downloaded and verified

```

[ ]: # Training loop
epochs = 10
for epoch in range(epochs):
    running_loss = 0.0
    for i, data in enumerate(trainloader):
        inputs, labels = data
        inputs = inputs.view(inputs.shape[0], -1)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f'Epoch {epoch + 1}, Loss: {running_loss / len(trainloader)}')

```

```

Epoch 1, Loss: 1.8407409086227418
Epoch 2, Loss: 1.691075718641281
Epoch 3, Loss: 1.6235661569833755
Epoch 4, Loss: 1.573109028339386
Epoch 5, Loss: 1.5269226109981537
Epoch 6, Loss: 1.486977286696434
Epoch 7, Loss: 1.446461871623993
Epoch 8, Loss: 1.412262617468834
Epoch 9, Loss: 1.379416198015213
Epoch 10, Loss: 1.348554687321186

```

```
[11]: # Test accuracy calculation
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images = images.view(images.size(0), -1)
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy of the network on the 10000 test images: {100 * correct /
↪total}%')
```

Accuracy of the network on the 10000 test images: 47.92%

2. Create a LeNet-5 network for MNIST using Pytorch. Report the training and test accuracies after 10 epochs.

```
[ ]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
```

```
[ ]: class LeNet5(nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, 5) # 1 input channel (grayscale), 6
↪output channels, 5x5 kernel size
        self.pool = nn.AvgPool2d(2, 2) # Average pooling with 2x2 kernel size
↪and stride 2
        self.conv2 = nn.Conv2d(6, 16, 5) # 6 input channels, 16 output
↪channels, 5x5 kernel size
        self.fc1 = nn.Linear(16 * 4 * 4, 120) # 16 channels, 4x4 image size
↪after pooling, 120 output features
        self.fc2 = nn.Linear(120, 84) # 120 input features, 84 output features
        self.fc3 = nn.Linear(84, 10) # 84 input features, 10 output features/
↪classes for MNIST

    def forward(self, x):
        x = self.pool(torch.tanh(self.conv1(x)))
        x = self.pool(torch.tanh(self.conv2(x)))
        x = x.view(-1, 16 * 4 * 4) # Flatten the 16x4x4 tensor into a 1D tensor
        x = torch.tanh(self.fc1(x))
        x = torch.tanh(self.fc2(x))
        x = self.fc3(x)
```

```
return x
```

```
[ ]: transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

trainset = torchvision.datasets.MNIST(root=r'D:\semi 5\IP&CV\Assignment_
↪3\data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=64, shuffle=True)

testset = torchvision.datasets.MNIST(root=r'D:\semi 5\IP&CV\Assignment 3\data',
↪train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=64, shuffle=False)
```

```
[ ]: model = LeNet5()
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

for epoch in range(10):
    running_loss = 0.0
    correct = 0
    total = 0
    for inputs, labels in trainloader:
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    train_accuracy = 100 * correct / total
    print(f'Epoch {epoch + 1}, Loss: {running_loss / len(trainloader):.4f},
↪Training Accuracy: {train_accuracy:.2f}%')

correct = 0
total = 0
with torch.no_grad():
    for inputs, labels in testloader:
        outputs = model(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
```



```
test_accuracy = 100 * correct / total
print(f'Test Accuracy: {test_accuracy:.2f}%')
```

```
Epoch 1, Loss: 1.5365, Training Accuracy: 55.22%
Epoch 2, Loss: 0.5226, Training Accuracy: 86.18%
Epoch 3, Loss: 0.3634, Training Accuracy: 89.91%
Epoch 4, Loss: 0.2954, Training Accuracy: 91.56%
Epoch 5, Loss: 0.2481, Training Accuracy: 92.93%
Epoch 6, Loss: 0.2115, Training Accuracy: 93.94%
Epoch 7, Loss: 0.1829, Training Accuracy: 94.75%
Epoch 8, Loss: 0.1607, Training Accuracy: 95.37%
Epoch 9, Loss: 0.1431, Training Accuracy: 95.85%
Epoch 10, Loss: 0.1289, Training Accuracy: 96.26%
Test Accuracy: 96.65%
```

**3. Based on the PyTorch tutorial on transfer learning get the pre-trained ResNet18 network trained on ImageNet15. classify hymenoptera dataset by**

- (a) fine tuning, and
- (b) using the network as a feature extractor.

```
[31]: import torch
import torchvision
import torchvision.transforms as transforms
from torchvision import datasets, models, transforms
import os
from torch import nn, optim
from torch.optim import lr_scheduler
import time
import copy
```

```
[33]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f'Using device: {device}')
```

```
model_conv = models.resnet18(pretrained=True)
```

```
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    'val': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
```

```

        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}

data_dir = r'D:\semi 5\IP&CV\Assignment 3\data\hymenoptera_data'
image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
    ↪data_transforms[x]) for x in ['train', 'val']}
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=4,
    ↪shuffle=True, num_workers=4) for x in ['train', 'val']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
class_names = image_datasets['train'].classes

```

Using device: cpu

```

[ ]: def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
    since = time.time()

    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    for epoch in range(num_epochs):
        print(f'Epoch {epoch}/{num_epochs - 1}')
        print('-' * 10)

        for phase in ['train', 'val']:
            if phase == 'train':
                model.train()  # Set model to training mode
            else:
                model.eval()   # Set model to evaluate mode

            running_loss = 0.0
            running_corrects = 0

            for inputs, labels in dataloaders[phase]:
                inputs = inputs.to(device)
                labels = labels.to(device)

                optimizer.zero_grad()

                with torch.set_grad_enabled(phase == 'train'): # Set gradient
    ↪calculation only for training phase
                    outputs = model(inputs)
                    _, preds = torch.max(outputs, 1)
                    loss = criterion(outputs, labels)

                    if phase == 'train': # Only perform backward and
    ↪optimization during training phase

```

```

        loss.backward()
        optimizer.step()

        running_loss += loss.item() * inputs.size(0)
        running_corrects += torch.sum(preds == labels.data)
    if phase == 'train':
        scheduler.step()

    epoch_loss = running_loss / dataset_sizes[phase]
    epoch_acc = running_corrects.double() / dataset_sizes[phase]

    print(f'{phase}- Loss: {epoch_loss:.4f} Accuracy: {epoch_acc:.4f}')

    if phase == 'val' and epoch_acc > best_acc: # Save the best model
        best_acc = epoch_acc
        best_model_wts = copy.deepcopy(model.state_dict())

    print()

    time_elapsed = time.time() - since
    print(f'Training complete in {time_elapsed // 60:.0f}m {time_elapsed % 60:.
↪0f}s')
    print(f'Best val Accuracy: {best_acc:.4f}')

    model.load_state_dict(best_model_wts)
    return model

```

#### (a) Fine tuning

```

[35]: model_ft = models.resnet18(pretrained=True)
      num_fts = model_ft.fc.in_features
      model_ft.fc = nn.Linear(num_fts, 2)

      model_ft = model_ft.to(device)
      criterion = nn.CrossEntropyLoss()

      optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

      exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)

      model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
↪num_epochs=25)

```

Epoch 0/24

-----

train- Loss: 0.5393 Accuracy: 0.7172

val- Loss: 0.2172 Accuracy: 0.9216

Epoch 1/24

-----

train- Loss: 0.4260 Accuracy: 0.8402

val- Loss: 0.4550 Accuracy: 0.8170

Epoch 2/24

-----

train- Loss: 0.5254 Accuracy: 0.7787

val- Loss: 0.3230 Accuracy: 0.9020

Epoch 3/24

-----

train- Loss: 0.6947 Accuracy: 0.7664

val- Loss: 0.2025 Accuracy: 0.9216

Epoch 4/24

-----

train- Loss: 0.6024 Accuracy: 0.7623

val- Loss: 0.2493 Accuracy: 0.8954

Epoch 5/24

-----

train- Loss: 0.4619 Accuracy: 0.8156

val- Loss: 0.3072 Accuracy: 0.8954

Epoch 6/24

-----

train- Loss: 0.8062 Accuracy: 0.6967

val- Loss: 0.5599 Accuracy: 0.8235

Epoch 7/24

-----

train- Loss: 0.4273 Accuracy: 0.8115

val- Loss: 0.3046 Accuracy: 0.9150

Epoch 8/24

-----

train- Loss: 0.4436 Accuracy: 0.8074

val- Loss: 0.2835 Accuracy: 0.9085

Epoch 9/24

-----

train- Loss: 0.3257 Accuracy: 0.8484

val- Loss: 0.2748 Accuracy: 0.9216

Epoch 10/24

-----

train- Loss: 0.2825 Accuracy: 0.8852

val- Loss: 0.2603 Accuracy: 0.9281

Epoch 11/24

-----

train- Loss: 0.2884 Accuracy: 0.8607

val- Loss: 0.2426 Accuracy: 0.9281

Epoch 12/24

-----

train- Loss: 0.2042 Accuracy: 0.9180

val- Loss: 0.2577 Accuracy: 0.9281

Epoch 13/24

-----

train- Loss: 0.2788 Accuracy: 0.8770

val- Loss: 0.2336 Accuracy: 0.9346

Epoch 14/24

-----

train- Loss: 0.2652 Accuracy: 0.8934

val- Loss: 0.2368 Accuracy: 0.9281

Epoch 15/24

-----

train- Loss: 0.2566 Accuracy: 0.9057

val- Loss: 0.2320 Accuracy: 0.9281

Epoch 16/24

-----

train- Loss: 0.3060 Accuracy: 0.8730

val- Loss: 0.2338 Accuracy: 0.9412

Epoch 17/24

-----

train- Loss: 0.3132 Accuracy: 0.8934

val- Loss: 0.2471 Accuracy: 0.9216

Epoch 18/24

-----

train- Loss: 0.2000 Accuracy: 0.9262

val- Loss: 0.2300 Accuracy: 0.9346

Epoch 19/24

-----

train- Loss: 0.2233 Accuracy: 0.9016

val- Loss: 0.2212 Accuracy: 0.9346

Epoch 20/24

```
-----  
train- Loss: 0.2678 Accuracy: 0.8689  
val- Loss: 0.2319 Accuracy: 0.9346
```

Epoch 21/24

```
-----  
train- Loss: 0.2537 Accuracy: 0.9057  
val- Loss: 0.2264 Accuracy: 0.9346
```

Epoch 22/24

```
-----  
train- Loss: 0.2131 Accuracy: 0.9262  
val- Loss: 0.2360 Accuracy: 0.9281
```

Epoch 23/24

```
-----  
train- Loss: 0.2720 Accuracy: 0.8730  
val- Loss: 0.2320 Accuracy: 0.9346
```

Epoch 24/24

```
-----  
train- Loss: 0.3468 Accuracy: 0.8402  
val- Loss: 0.2212 Accuracy: 0.9412
```

Training complete in 13m 11s  
Best val Accuracy: 0.9412

### (b) Network as a feature extractor

```
[36]: model_conv = torchvision.models.resnet18(pretrained=True)  
      for param in model_conv.parameters():  
          param.requires_grad = False  
  
      num_fters = model_conv.fc.in_features  
      model_conv.fc = nn.Linear(num_fters, 2)  
  
      model_conv = model_conv.to(device)  
      criterion = nn.CrossEntropyLoss()  
  
      optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=0.001, momentum=0.9)  
  
      exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=7, gamma=0.1)  
  
      model_conv = train_model(model_conv, criterion, optimizer_conv,   
                               ↪exp_lr_scheduler, num_epochs=25)
```

Epoch 0/24

```
-----  
train- Loss: 0.5542 Accuracy: 0.7172
```

val- Loss: 0.2496 Accuracy: 0.9085

Epoch 1/24

-----

train- Loss: 0.4786 Accuracy: 0.7541

val- Loss: 0.1975 Accuracy: 0.9216

Epoch 2/24

-----

train- Loss: 0.4831 Accuracy: 0.7910

val- Loss: 0.5865 Accuracy: 0.7778

Epoch 3/24

-----

train- Loss: 0.4871 Accuracy: 0.7664

val- Loss: 0.2344 Accuracy: 0.9150

Epoch 4/24

-----

train- Loss: 0.5491 Accuracy: 0.7541

val- Loss: 0.2346 Accuracy: 0.9216

Epoch 5/24

-----

train- Loss: 0.5867 Accuracy: 0.7541

val- Loss: 0.3169 Accuracy: 0.8889

Epoch 6/24

-----

train- Loss: 0.3908 Accuracy: 0.8197

val- Loss: 0.2522 Accuracy: 0.9085

Epoch 7/24

-----

train- Loss: 0.3870 Accuracy: 0.8238

val- Loss: 0.2197 Accuracy: 0.9216

Epoch 8/24

-----

train- Loss: 0.3586 Accuracy: 0.8443

val- Loss: 0.1810 Accuracy: 0.9412

Epoch 9/24

-----

train- Loss: 0.3468 Accuracy: 0.8197

val- Loss: 0.1811 Accuracy: 0.9477

Epoch 10/24

-----  
train- Loss: 0.3483 Accuracy: 0.8443  
val- Loss: 0.1770 Accuracy: 0.9281

Epoch 11/24

-----  
train- Loss: 0.3391 Accuracy: 0.8443  
val- Loss: 0.1689 Accuracy: 0.9412

Epoch 12/24

-----  
train- Loss: 0.3749 Accuracy: 0.8197  
val- Loss: 0.1942 Accuracy: 0.9412

Epoch 13/24

-----  
train- Loss: 0.3013 Accuracy: 0.8607  
val- Loss: 0.1868 Accuracy: 0.9412

Epoch 14/24

-----  
train- Loss: 0.2513 Accuracy: 0.8934  
val- Loss: 0.1828 Accuracy: 0.9412

Epoch 15/24

-----  
train- Loss: 0.2469 Accuracy: 0.9057  
val- Loss: 0.1729 Accuracy: 0.9477

Epoch 16/24

-----  
train- Loss: 0.2748 Accuracy: 0.9016  
val- Loss: 0.2437 Accuracy: 0.9216

Epoch 17/24

-----  
train- Loss: 0.2543 Accuracy: 0.9016  
val- Loss: 0.1866 Accuracy: 0.9346

Epoch 18/24

-----  
train- Loss: 0.3329 Accuracy: 0.8525  
val- Loss: 0.1675 Accuracy: 0.9412

Epoch 19/24

-----  
train- Loss: 0.3754 Accuracy: 0.8484  
val- Loss: 0.1705 Accuracy: 0.9412



Epoch 20/24

-----

train- Loss: 0.3203 Accuracy: 0.8525

val- Loss: 0.2097 Accuracy: 0.9216

Epoch 21/24

-----

train- Loss: 0.2706 Accuracy: 0.8811

val- Loss: 0.1968 Accuracy: 0.9281

Epoch 22/24

-----

train- Loss: 0.2644 Accuracy: 0.8934

val- Loss: 0.2084 Accuracy: 0.9216

Epoch 23/24

-----

train- Loss: 0.3636 Accuracy: 0.8484

val- Loss: 0.1918 Accuracy: 0.9281

Epoch 24/24

-----

train- Loss: 0.2513 Accuracy: 0.8975

val- Loss: 0.2239 Accuracy: 0.9281

Training complete in 8m 4s

Best val Accuracy: 0.9477